


Performance Tips for container applications

How to test and check performance issues



Schedule

- Infrastructure and Application
- What look?
- How to look?
- Optimization tradeoffs



Infrastructure - Concerns

- Isolated
- One container application doesn't affect other
- You don't need to worry about where your container will run
- You can reproduce the same behavior locally



Infrastructure - Concerns

- Your application will share resources with other containers?
 - **Pros**
 - Better use of your physical resources (CPU, Memory, etc)
 - Your application it is not using CPU ALL the time
 - Make sense have more than one container in the same host
 - Save money
 - **Cons**
 - The physical resources available changes over time
 - Application impacts each other
 - Spent more money




Infrastructure - Concerns

- All containers in the same host are yours?
 - If yes
 - The control about sharing resources belongs to the team responsible by the applications
 - The development team needs to know about infrastructure
 - The development team could tune the applications and hardware
 - Could have waste of resources (Ex: batch operations)
 - If not
 - The control about sharing resources belongs to the infrastructure team
 - The development team has fewer options to tune application
 - Possible no waste of resources

Infrastructure - Concerns

- The host will share containers with other applications?
 - Data base
 - Infrastructure resources
 - Load Balancers
 - HAProxy
 - Nginx / Apache
 - Cache Layer
 - Logger App (ex: FluentD)



Infrastructure - Concerns

- The environment will put some limits?
 - Memory
 - CPU usage
 - File Descriptor
 - Network Bandwidth
 - Disk I/O operations



Application - Concerns

- Your application will run inside a Virtual Machine / Server ?
 - Java with JVM (Application Server: Tomcat, Jetty, etc)
 - .NET Kestrel Web Server
- Your container runs more than one application?



What look?

- Memory
- CPU usage
- File Descriptor
 - Sockets
 - Files
- Network Bandwidth
- Disk I/O operations



How to look? - Docker run command

```
docker run -it --ulimit nofile=<file_descriptor_limit>  
--memory <memory_limit> --cpus <cpu_limit> --name <app_name>  
<app_image>
```

- **Hardware/Infrastructure Limits:**

- <file_descriptor_limit>
- <memory_limit>
- <cpu_limit>

How to look? - Minimum requirements

- Run your application without limits
- Performance Test with status URL (do nothing)
- Performance Test with service URL (do the thing)



How to look? - Example

- Java Application
- Endpoints
 - GET /status : do nothing, return HTTP 200
 - POST /sort : Sort N lists of integer numbers internally using async for each list and with thread sleep with random milliseconds



{ REST:API }

How to look? - Example

```
curl -v -X POST http://localhost:8080/sort 'content-type: application/json' \  
-d '{  
  "intList": [{  
    "list": [8, 9, 4, 12, 56, 78, 90, 56, 34, 74]  
  }, {  
    "list": [8, 9, 4, 12, 56, 78, 90, 56, 34, 74]  
  }, {  
    "list": [8, 9, 4, 12, 56, 78, 90, 56, 34, 74]  
  }, {  
    "list": [8, 9, 4, 12, 56, 78, 90, 56, 34, 74]  
  }, {  
    "list": [8, 9, 4, 12, 56, 78, 90, 56, 34, 74]  
  }]  
}'
```

How to look? - Example

Run Application

```
docker run -it -p 8080:8080 --name demo demo:1.0.0
```

Docker Status

```
docker stats demo
```

Open File Descriptors

```
docker exec -it demo /bin/bash  
apt-get update  
apt-get install procps -y  
ps -efl  
watch 'ls -altr /proc/6/fd | wc -l'
```

How to look? - Example

Performance Test (Ex: Apache Benchmark Tool)

```
ab -s 60 -c 10 -n 20 http://localhost:8080/status
```

- **-c : Concurrent Users**
- **-n : Number of iterations**
- **-s : timeout (seconds)**

ApacheBench



Windows PowerShell

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O	BLOCK I/O
7d6078cf18be	demo	0.05%	287.1MiB / 1.934GiB	14.50%	9.98kB / 9.73kB	3.26MB / 0B

Windows PowerShell

```
Every 2.0s: ls -altr /proc/6/fd | wc -l                                7d6078cf18be: T
40
```

Windows PowerShell

```
2018-10-02 13:36:34.867 INFO 6 --- [          main] c.e.m.MicroserviceDemoApplication : Started MicroserviceDemoApplication
unning for 2.908)
2018-10-02 13:45:19.217 INFO 6 --- [nio-8080-exec-1] o.a.c.c.C.[Tomcat].[localhost].[/] : Initializing Spring FrameworkServlet
2018-10-02 13:45:19.217 INFO 6 --- [nio-8080-exec-1] o.s.web.servlet.DispatcherServlet : FrameworkServlet 'dispatcherServlet'
d2018-10-02 13:45:19.240 INFO 6 --- [nio-8080-exec-1] o.s.web.servlet.DispatcherServlet : FrameworkServlet 'dispatcherServlet'
eted in 22 ms
```

Select Windows PowerShell

	min	mean[+/-sd]	median	max
Connect:	0	0 0.5	0	1
Processing:	22	57 40.8	53	216
Waiting:	14	46 36.5	42	188
Total:	22	57 41.1	53	217

Percentage of the requests served within a certain time (ms)

50%	53
66%	61
75%	63
80%	72
90%	77
95%	217
98%	217
99%	217
100%	217 (longest request)

PS C:\Users\felipe.lino>

How to look? - Example

Performance Test (Ex: Apache Benchmark Tool)

```
ab -p request.json -T application/json -H 'accept:  
application/json' -c 10 -n 200 http://localhost:8080/sort
```

Windows PowerShell

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O	BLOCK I/O
7d6078cf18be	demo	0.07%	301MiB / 1.934GiB	15.20%	187kB / 193kB	3.47MB / 0B

Windows PowerShell

Every 2.0s: ls -altr /proc/6/fd | wc -l 7d6078cf18be:

40

Windows PowerShell

work.aop.interceptor.AsyncExecutionInterceptor\$\$Lambda\$335/1679947981@41b26123] with root cause

java.util.concurrent.RejectedExecutionException: Task org.springframework.util.concurrent.ListenableFutureTask@2971e16e rejected from ThreadPoolExecutor@5961e63a[Running, pool size = 2, active threads = 2, queued tasks = 3, completed tasks = 60]

at java.util.concurrent.ThreadPoolExecutor\$AbortPolicy.rejectedExecution(ThreadPoolExecutor.java:2063) ~[na:1.8.0_181]

at java.util.concurrent.ThreadPoolExecutor.reject(ThreadPoolExecutor.java:830) [na:1.8.0_181]

at java.util.concurrent.ThreadPoolExecutor.execute(ThreadPoolExecutor.java:1379) [na:1.8.0_181]

Windows PowerShell

Connection Times (ms)

	min	mean[+/-sd]	median	max
Connect:	0	0 0.4	0	1
Processing:	15	93 40.7	87	231
Waiting:	12	85 37.4	81	230
Total:	15	93 40.7	87	231

Percentage of the requests served within a certain time (ms)

50%	87
66%	103
75%	114
80%	123
90%	151
95%	175
98%	214
99%	223
100%	231 (longest request)

PS C:\Users\felipe.lino\github\container-optimizer-demo> ab -p request.json -T application/json -H 'accept: application/json' -c 10

How to look? - Maximum requirements

- Run your application applying infrastructure limits
- Run your application applying server (JVM) parameters
- Performance Test with service URL (do the thing)



How to look? - Example

Run Application

```
docker run -it --ulimit nofile=128 --memory 1000MB --cpus 0.5 -e  
"JAVA_OPTS=-Xms312m -Xmx750m -DcorePoolSize=100 -DmaxPoolSize=120  
-DqueueCapacity=300" -p 8080:8080 --name demo demo:1.0.0
```

Infrastructure limits:

- --cpus
- --memory
- --ulimit

Applications config:

- Threads: corePoolSize (2), maxPoolSize (2), queueCapacity(3)
- Xms / Xmx
- etc.

How to look? Before start again

Windows PowerShell

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O	BLOCK I/O
17f52abc18ff	demo	0.05%	300.1MiB / 1000MiB	30.01%	1.04kB / 0B	442kB / 0B

Windows PowerShell

```
Every 2.0s: ls -altr /proc/6/fd | wc -l 17f52abc18ff:
39
```

- Memory Limit is 1GB
- You start to usage around 300MB because of JAVA_OPTS with *-Xms312m*

Windows PowerShell

CONTAINER ID	NAME	CPU %	MEM USAGE / LIMIT	MEM %	NET I/O	BLOCK I/O
c4a0020e315e	demo	0.36%	288.6MiB / 1000MiB	28.86%	162kB / 132kB	303kB / 0B

Windows PowerShell

```
Every 2.0s: ls -altr /proc/6/fd | wc -l c4a0020e315e:
40
```

Windows PowerShell

```
2018-10-02 14:21:22.554 INFO 6 --- [CustomThread-43] c.e.m.service.SortServiceImpl : Sleeping for 98 ms
2018-10-02 14:21:22.554 INFO 6 --- [CustomThread-86] c.e.m.service.SortServiceImpl : Sleeping for 82 ms
2018-10-02 14:21:22.554 INFO 6 --- [CustomThread-87] c.e.m.service.SortServiceImpl : Sleeping for 36 ms
2018-10-02 14:21:22.558 INFO 6 --- [CustomThread-94] c.e.m.service.SortServiceImpl : InputList: [8, 9, 4, 12, 56, 78, 9
2018-10-02 14:21:22.558 INFO 6 --- [CustomThread-94] c.e.m.service.SortServiceImpl : Sleeping for 31 ms
2018-10-02 14:21:22.568 INFO 6 --- [CustomThread-33] c.e.m.service.SortServiceImpl : OutputList: [4, 8, 9, 12, 34, 56,
2018-10-02 14:21:22.589 INFO 6 --- [CustomThread-40] c.e.m.service.SortServiceImpl : OutputList: [4, 8, 9, 12, 34, 56,
```


Windows PowerShell

```
Connection Times (ms)
      min      mean[+/-sd]    median    max
Connect:      0       0   0.4         0      1
Processing:  92    191 105.3      184   1245
Waiting:     80    183  95.0      178  1038
Total:       92    192 105.3      184   1245

Percentage of the requests served within a certain time (ms)
 50%    184
 66%    196
 75%    205
 80%    216
 90%    285
 95%    322
 98%    504
 99%    524
100%   1245 (longest request)
PS C:\Users\felipe.lino\github\container-optimizer-demo>
```

Optimizations tradeoffs

- **Change memory limits**
- **Change CPU limits**
- **File Descriptors and Sockets**
- **Number of containers x Threads (workers) inside application**
- **Specific configurations for your Server/JVM**
 - `server.undertow.io-threads`
 - `KestrelServerOptions.ThreadCount`



Optimizations tradeoffs

- **IO Usage**
- **Network Usage**
- **Related dependencies**
 - **APIs**
 - **Database**
 - **File System**

Any doubts?



Source:

<https://github.com/felipelino/container-optimizer-demo>

Contacts:

Felipe Lino (felipelino44@gmail.com)